

Handling Missing Data in ETL

1. What are the most common reasons for missing data in ETL pipelines?

Ans.

- Source system issues → Data not captured due to faulty sensors, human error, or incomplete forms.
- Schema changes upstream → Column names or structures change, breaking extraction.
- Integration errors → API failures, connection timeouts, or mismatched formats during extraction.
- Data corruption → Loss during transfer because of network latency or hardware failures.
- Workflow bottlenecks → Jobs fail mid-process due to resource limits or poor orchestration.
- Business rules → Certain fields intentionally left blank (e.g., optional attributes).

2. Why is blindly deleting rows with missing values considered a bad practice in ETL?

Ans.

- Loss of valuable information → Entire records are discarded even if most fields are valid.
- Bias in analysis → If missingness is systematic (e.g., certain regions or customer groups), deletion skews results.
- Reduced sample size → Smaller datasets weaken statistical power and model accuracy.
- Business impact → Important insights (like why data is missing) are lost, which could itself be meaningful.

3. Q3. Explain the difference between:

- Listwise deletion
- Column deletion

Also mention one scenario where each is appropriate. Ans.

Listwise Deletion

- Removes entire rows if any required field is missing.
- Preserves schema but reduces dataset size.
- Scenario: Customer survey data where multiple answers are missing → better to drop the whole record to avoid incomplete profiles.

Column Deletion

- Removes entire columns if they have too many missing values.
- Preserves row count but loses that variable.

- Scenario: A dataset where “Middle Name” is missing for 95% of customers → dropping the column avoids clutter without harming analysis.

4. Why is median imputation preferred over mean imputation for skewed data such as income?

Ans.

- Mean imputation → Sensitive to extreme values (outliers). In skewed distributions like income, a few very high salaries can inflate the mean, making it unrepresentative.
- Median imputation → More robust because the median is the middle value, unaffected by outliers. It better reflects the “typical” case in skewed datasets.
- Example: If most incomes are around 50,000 but a few are above 1,000,000, the mean will be distorted, while the median stays close to the majority.

5. What is forward fill and in what type of dataset is it most useful?

Ans.

- ForwardFill(FFILL):

A technique where missing values are replaced with the last valid observation carried forward.
- Usefulness:
 - Best suited for time-series or sequential datasets (e.g., monthly sales, stock prices, sensor readings).
 - Assumes continuity—the last known value is a reasonable proxy until a new one appears.
- Example:

If sales data for March is missing but February had 12,000, forward fill assigns March = 12,000 until April's actual figure arrives.

6. Why should flagging missing values be done before imputation in an ETL workflow?

Ans.

- Preserves information about missingness → The fact that a value is missing can itself carry business meaning (e.g., customers not disclosing income).
- Avoids loss of transparency → Once imputation is applied, you can't distinguish between original and filled values unless flagged beforehand.
- Supports better analysis → Analysts can study patterns of missingness separately (e.g., which regions or groups have more missing data).
- Improves model accuracy → Machine learning models can use the flags as an additional feature, helping them account for missingness bias.

7. Consider a scenario where income is missing for many customers. How can this missingness itself provide business insights?

Ans.

- Customer reluctance → Missing income may indicate customers are unwilling to disclose sensitive financial details, signaling trust or privacy concerns.
- Segment behavior → Certain demographics or regions may have higher missingness, revealing cultural or socio-economic differences.
- Product targeting → If high-value customers avoid sharing income, it may suggest they don't see relevance in providing it, guiding product design or survey strategy.
- Operational gaps → Consistent missingness could highlight flaws in data collection

